



---

Welcome to Hive



---

# Hive - Introduction

---

- Data warehouse infrastructure tool
- Process structured data in Hadoop
- Resides on top of Hadoop
- Makes data churning easy
- Provides SQL like queries

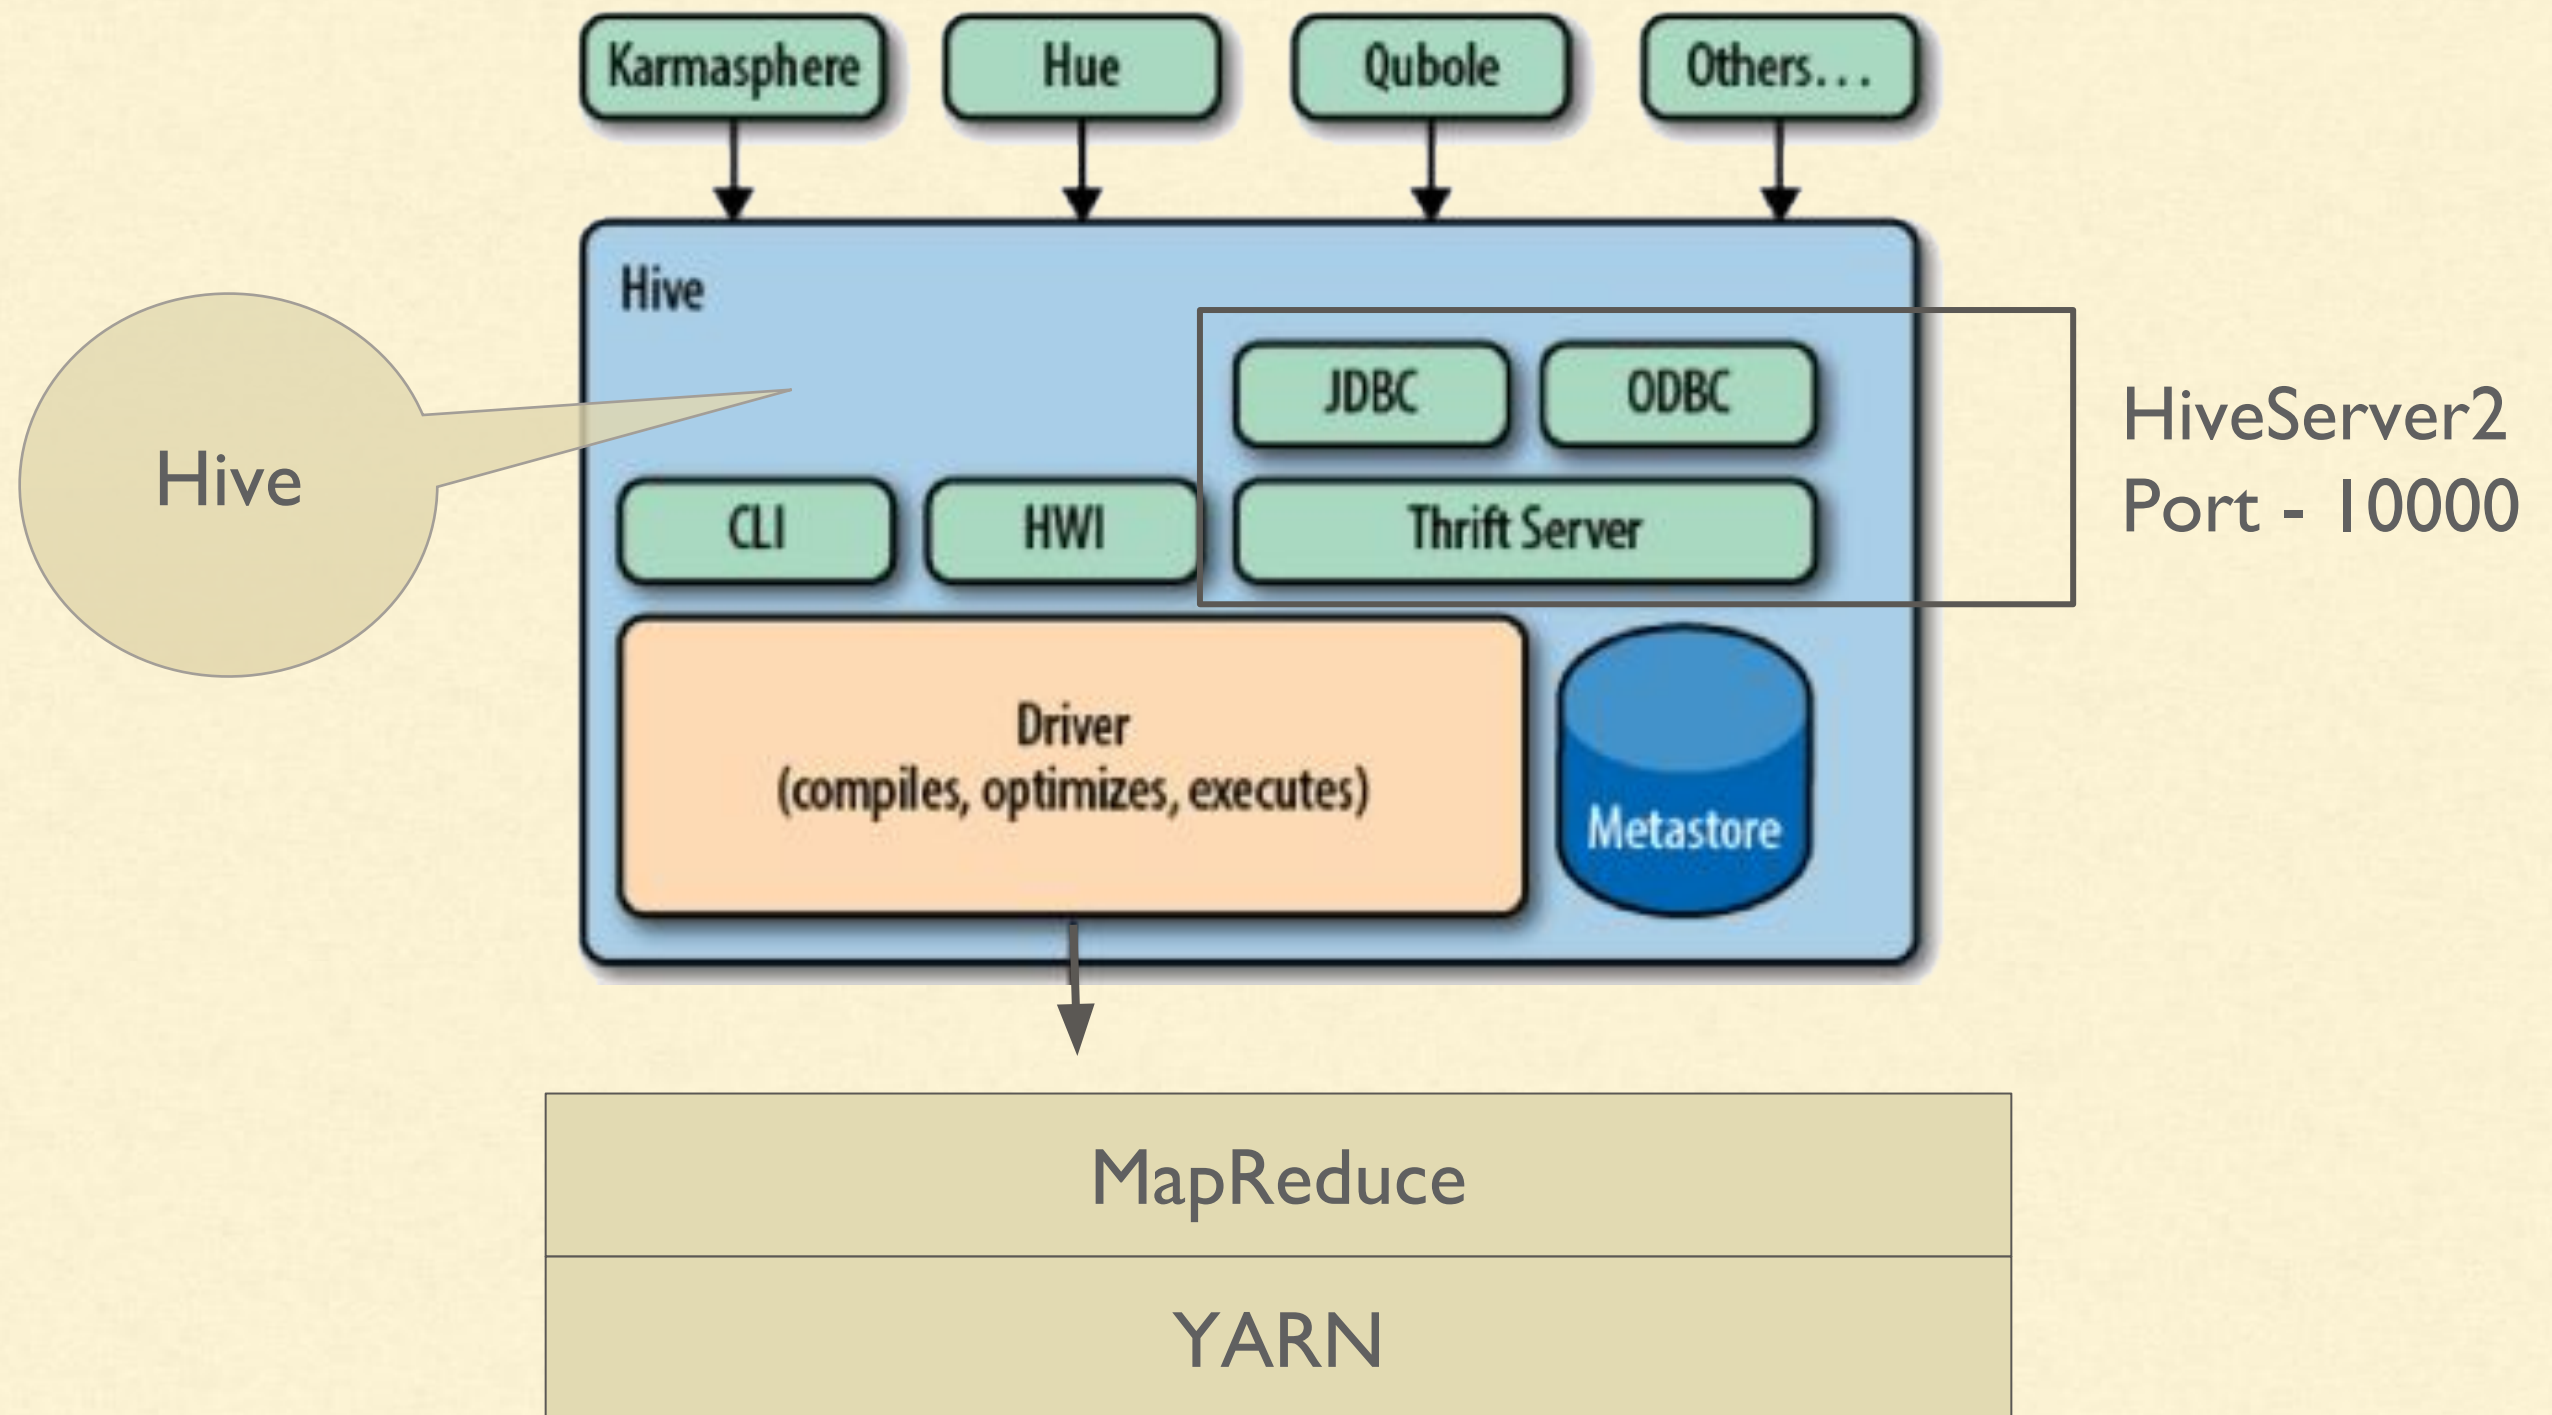
---

# Why Do We Need Hive?

---

- Developers face problem in writing MapReduce logic
- How to port existing
  - relational databases
  - SQL infrastructure with Hadoop?
- End users are familiar with SQL queries than MapReduce and Pig
- Hive's SQL-like query language makes data churning easy

# Hive - Components





---

# Hive - Limitations

---

- Does not provide row level updates (earlier versions)
- Not suitable for OLTP
  - Queries have higher latency
  - Start-up overhead for MapReduce jobs
- Best when large dataset is maintained and mined

---

# Hive - Data Types - Numeric

---

- TINYINT (1-byte signed integer)
- SMALLINT (2-byte signed integer)
- INT (4-byte signed integer)
- BIGINT (8-byte signed integer)
- FLOAT (4-byte single precision floating point number)
- DOUBLE (8-byte double precision floating point number)
- DECIMAL (User defined precisions)

---

# Hive - Data Types - Date/Time

---

- TIMESTAMP ( Hive version > 0.8.0 )
- DATE ( Hive version > 0.12.0 ) - YYYY-MM-DD

---

# Hive - Data Types - String

---

- STRING
- VARCHAR ( Hive version  $> 0.12.0$  )
- CHAR ( Hive version  $> 0.13.0$  )



---

# Hive - Data Types - Misc

---

- BOOLEAN
- BINARY ( Hive version > 0.8.0 )

---

# Hive - Data Types - Complex

---

arrays: ARRAY<data\_type>

maps: MAP<primitive\_type, data\_type>

structs: STRUCT<col\_name : data\_type [COMMENT col\_comment], ...>

union: UNIONTYPE<data\_type, data\_type, ...> ( Hive version > 0.7.0 )

---

# Hive - Data Types - Example

---

```
CREATE TABLE employees(  
  name STRING,  
  salary FLOAT,  
  subordinates ARRAY<STRING>,  
  deductions MAP<STRING, FLOAT>,  
  address STRUCT<street:STRING,  
    city:STRING,  
    state:STRING,  
    zip:INT>,  
  auth UNION<fbid:INT, gid:INT, email:STRING>  
)
```

---

# Hive - Data Types - Example

---

```
CREATE TABLE employees(  
  name STRING,  
  salary FLOAT,  
  subordinates ARRAY<STRING>,  
  deductions MAP<STRING, FLOAT>,  
  address STRUCT<street:STRING,  
    city:STRING,  
    state:STRING,  
    zip:INT>,  
  auth UNION<fbid:INT, gid:INT, email:STRING>  
)
```



“John”



# Hive - Data Types - Example

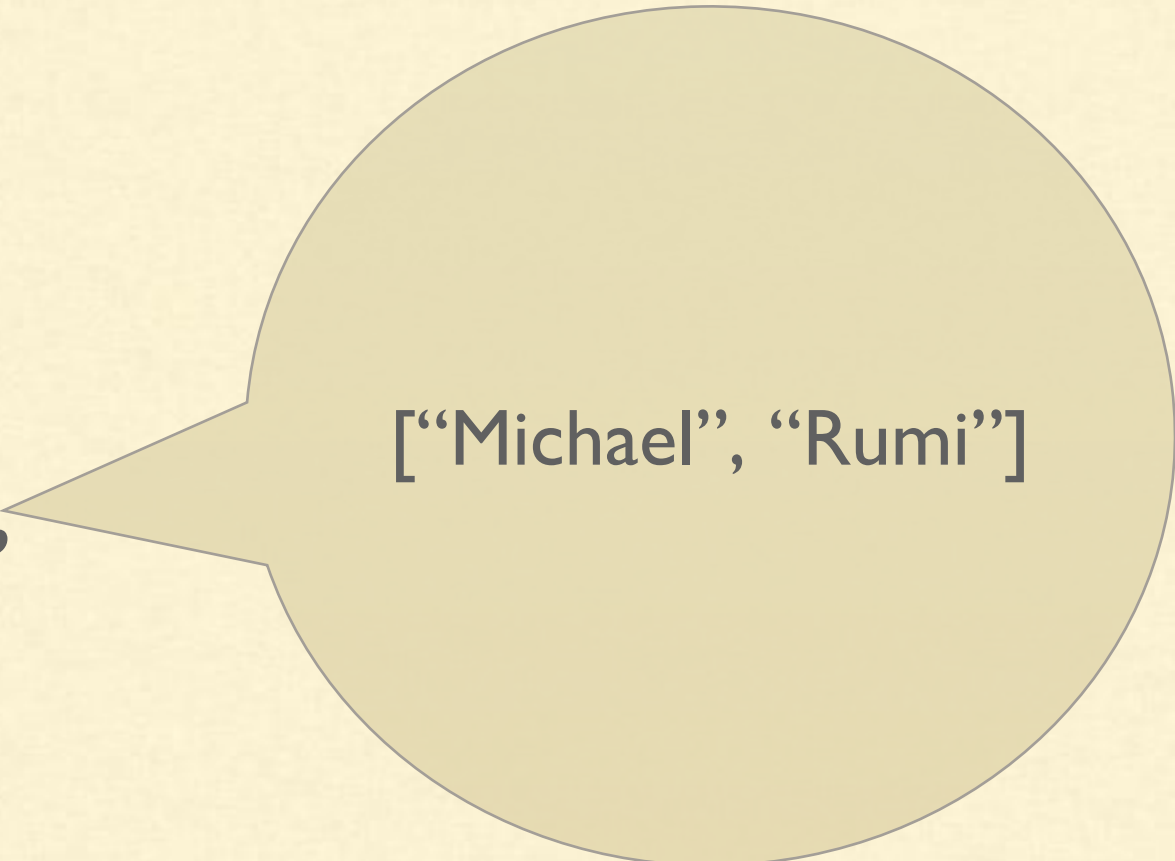
```
CREATE TABLE employees(  
  name STRING,  
  salary FLOAT,  
  subordinates ARRAY<STRING>,  
  deductions MAP<STRING, FLOAT>,  
  address STRUCT<street:STRING,  
    city:STRING,  
    state:STRING,  
    zip:INT>,  
  auth UNION<fbid:INT, gid:INT, email:STRING>  
)
```



40000.00

# Hive - Data Types - Example

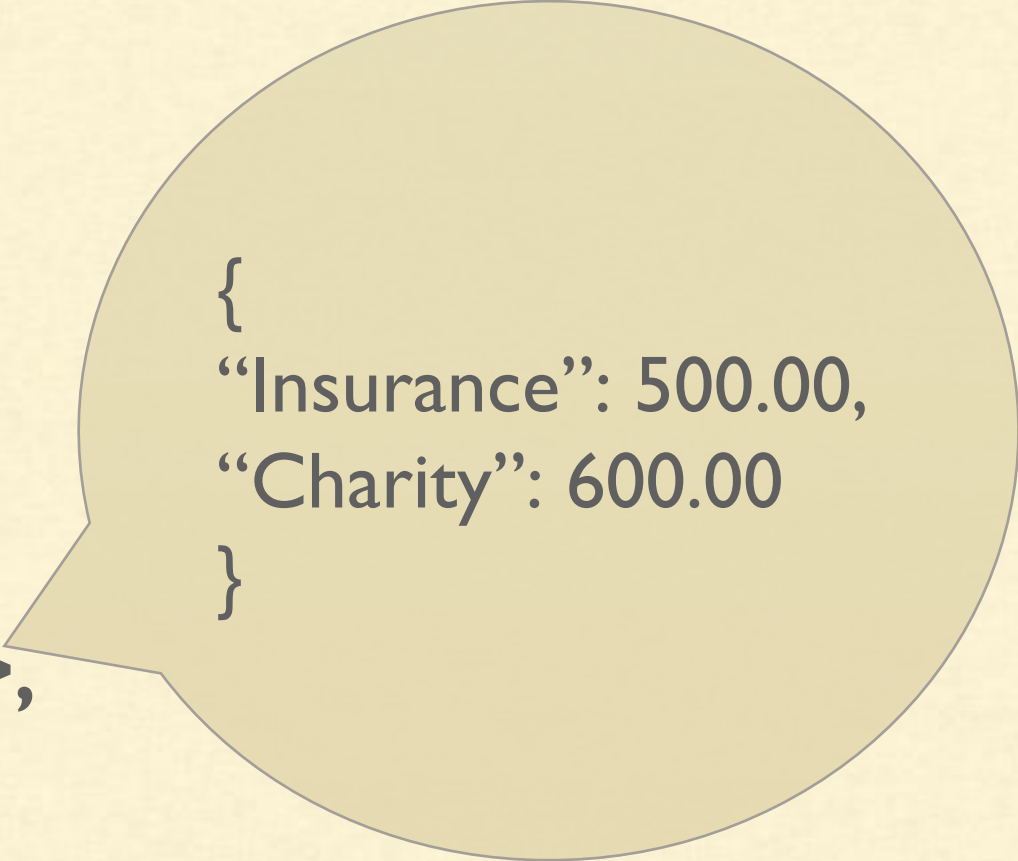
```
CREATE TABLE employees(  
  name STRING,  
  salary FLOAT,  
  subordinates ARRAY<STRING>,  
  deductions MAP<STRING, FLOAT>,  
  address STRUCT<street:STRING,  
    city:STRING,  
    state:STRING,  
    zip:INT>,  
  auth UNION<fbid:INT, gid:INT, email:STRING>  
)
```



[“Michael”, “Rumi”]

# Hive - Data Types - Example

```
CREATE TABLE employees(  
  name STRING,  
  salary FLOAT,  
  subordinates ARRAY<STRING>,  
  deductions MAP<STRING, FLOAT>,  
  address STRUCT<street:STRING,  
    city:STRING,  
    state:STRING,  
    zip:INT>,  
  auth UNION<fbid:INT, gid:INT, email:STRING>  
)
```



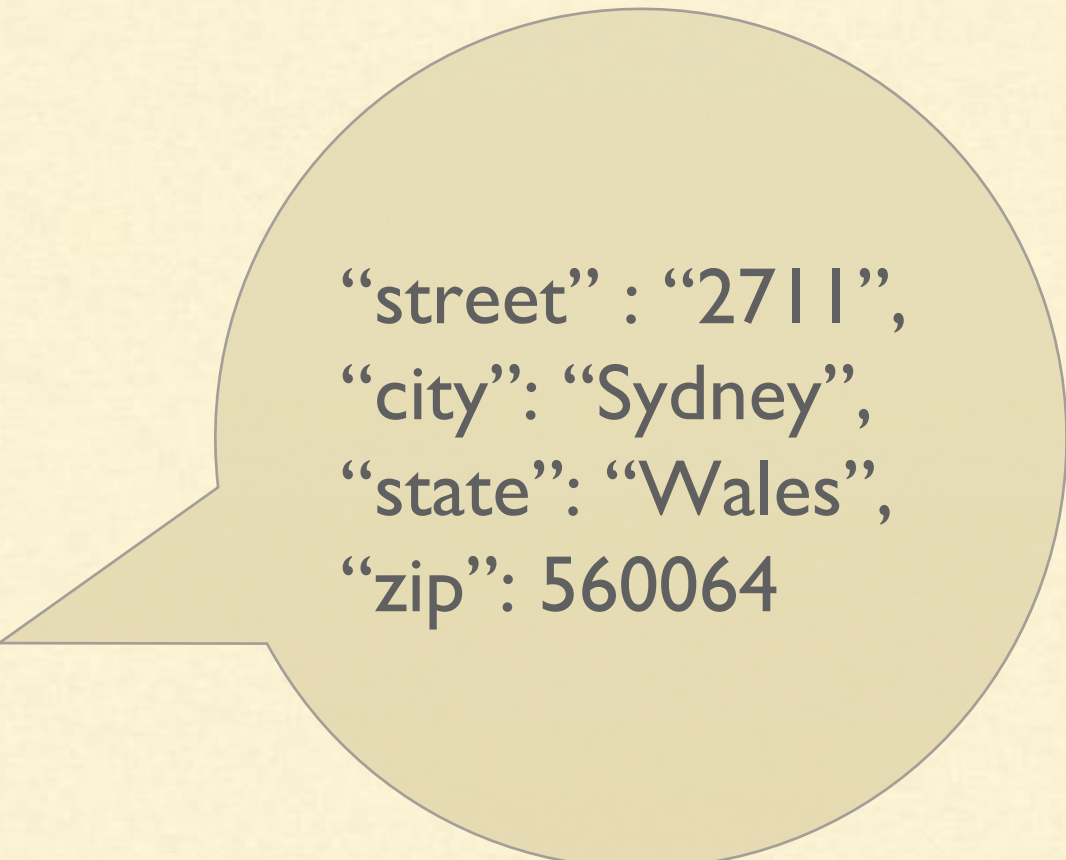
```
{  
  "Insurance": 500.00,  
  "Charity": 600.00  
}
```

---

# Hive - Data Types - Example

---

```
CREATE TABLE employees(  
  name STRING,  
  salary FLOAT,  
  subordinates ARRAY<STRING>,  
  deductions MAP<STRING, FLOAT>,  
  address STRUCT<street:STRING,  
    city:STRING,  
    state:STRING,  
    zip:INT>,  
  auth UNION<fbid:INT, gid:INT, email:STRING>  
)
```



“street” : “2711”,  
“city”: “Sydney”,  
“state”: “Wales”,  
“zip”: 560064



---

# Hive - Data Types - Example

---

```
CREATE TABLE employees(  
  name STRING,  
  salary FLOAT,  
  subordinates ARRAY<STRING>,  
  deductions MAP<STRING, FLOAT>,  
  address STRUCT<street:STRING,  
    city:STRING,  
    state:STRING,  
    zip:INT>,  
  auth UNION<fbid:INT, gid:INT, email:STRING>  
)
```



“fbid”:168478292

---

# Hive - Metastore

---

- Stores the metadata of tables into a relational database
- Metadata includes
  - Details of databases and tables
  - Table definitions: name of table, columns, partitions etc.

---

# Hive - Warehouse

---

- Hive tables are stored in the Hive warehouse directory
- `/apps/hive/warehouse` on HDFS
- At the location specified in the table definition

---

# Hive - Getting Started - Command Line

---

- Login to CloudeXLab Linux console
- Type “*hive*” to access hive shell
- By default database named “default” will be selected as current db for the current session
- Type “*SHOW DATABASES*” to see list of all databases



---

# Hive - Getting Started - Command Line

---

- “*SHOW TABLES*” will list tables in current selected database which is “default” database.
- Create your own database with your login name
- *CREATE DATABASE abhinav9884;*
- *DESCRIBE DATABASE abhinav9884;*
- *DROP DATABASE abhinav9884;*

---

# Hive - Getting Started - Command Line

---

- *CREATE DATABASE abhinav9884;*
- *USE abhinav9884;*
- *CREATE TABLE x (a INT);*

---

# Hive - Getting Started - Hue

---

- Login to Hue
- Click on “Query Editors” and select “Hive”
- Select your database (abhinav9984) from the list
- *SELECT \* FROM x;*
- *DESCRIBE x;*
- *DESCRIBE FORMATTED x;*

---

# Hive - Tables

---

- Managed tables
- External tables



---

# Hive - Managed Tables

---

- Aka Internal
- Lifecycle managed by Hive
- Data is stored in the warehouse directory
- Dropping the table deletes data from warehouse

---

# Hive - External Tables

---

- The lifecycle is not managed by Hive
- Hive assumes that it does not own the data
- Dropping the table does not delete the underlying data
- Metadata will be deleted

---

# Hive - Managed Tables - Example

---

```
CREATE TABLE nyse(  
    exchange1 STRING,  
    symbol1 STRING,  
    ymd STRING,  
    price_open FLOAT,  
    price_high FLOAT,  
    price_low FLOAT,  
    price_close FLOAT,  
    volume INT,  
    price_adj_close FLOAT  
)
```

```
ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t';
```

```
DESCRIBE nyse;  
DESCRIBE FORMATTED nyse;
```

---

# Hive - Loading Data - From Local Directory

---

- *hadoop fs -copyToLocal /data/NYSE\_daily*
- Launch Hive
- *use yourdatabase;*
- *load data local inpath 'NYSE\_daily' overwrite into table nyse;*
- Copies the data from local file system to warehouse



---

# Hive - Loading Data - From HDFS

---

```
CREATE TABLE nyse_hdfs(  
    exchange1 STRING,  
    symbol1 STRING,  
    ymd STRING,  
    price_open FLOAT,  
    price_high FLOAT,  
    price_low FLOAT,  
    price_close FLOAT,  
    volume INT,  
    price_adj_close FLOAT  
)  
ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t';
```

---

# Hive - Loading Data - From HDFS

---

- Copy /data/NYSE\_daily to your home directory in HDFS
- *load data inpath 'hdfs:///user/abhinav9884/NYSE\_daily' overwrite into table nyse\_hdfs;*
- Moves the data from specified location to warehouse
- Check if NYSE\_daily is in your home directory in HDFS

---

# Hive - External Tables

---

```
CREATE EXTERNAL TABLE nyse_external (  
    exchange1 STRING,  
    symbol1 STRING,  
    ymd STRING,  
    price_open FLOAT,  
    price_high FLOAT,  
    price_low FLOAT,  
    price_close FLOAT,  
    volume INT,  
    price_adj_close FLOAT  
)  
ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t'  
LOCATION '/user/abhinav9884/NYSE_daily';  
  
describe formatted nyse_external;
```

---

# Hive - S3 Based External Table

---

```
create external table miniwikistats (  
    projcode string,  
    pagename string,  
    pageviews int,  
    bytes int)  
partitioned by(dt string)  
row format delimited fields terminated by ' '  
lines terminated by '\n'  
location 's3n://paid/default-datasets/miniwikistats/';
```



---

# Hive - Select Statements

---

- Select all columns

*SELECT \* FROM nyse;*

- Select only required columns

*SELECT exchange1, symbol1 FROM nyse;*

---

# Hive - Aggregations

---

- Find average opening price for each stock

*SELECT symbol1, AVG(price\_open) AS avg\_price FROM nyse  
GROUP BY symbol1;*

- To improve performance set top-level aggregation in map phase

*SET hive.map.aggr=true;*

---

# Hive - Saving Data

---

- In local file system

*insert overwrite **local** directory '/home/abhinav9884/onlycmc'  
select \* from nyse where symbol = 'CMC';*

- In HDFS

*insert overwrite directory 'onlycmc' select \* from nyse where  
symbol = 'CMC';*

---

# Hive - Tables - DDL - ALTER

---

- Rename a table

*ALTER TABLE x RENAME TO x1;*

- Change datatype of column

*ALTER TABLE x1 CHANGE a a FLOAT;*

- Add columns in existing table

*ALTER TABLE x1 ADD COLUMNS (b FLOAT, c INT);*



---

# Hive - Partitions

---

#First name, Department, Year of joining

Mark, Engineering, 2012

Jon, HR, 2012

Monica, Finance, 2015

Steve, Engineering, 2012

Michael, Marketing, 2015

---

# Hive - Partitions - Hands-on

---

- Data is located at /data/bdhs/employees/ on HDFS
- Copy data to your home directory in HDFS

*hadoop fs -cp /data/bdhs/employees .*

- Create table

```
CREATE TABLE employees(  
name STRING,  
department STRING,  
somedate DATE  
)  
PARTITIONED BY(year STRING)  
ROW FORMAT DELIMITED FIELDS TERMINATED BY ',';
```

---

# Hive - Partitions - Hands-on

---

- Load dataset 2012.csv

*load data inpath 'hdfs:///user/sandeepgiri9034/employees/2012.csv' into table employees partition (year=2012);*

- Load dataset 2015.csv

*load data inpath 'hdfs:///user/sandeepgiri9034/employees/2015.csv' into table employees partition (year=2015);*

- *SHOW PARTITIONS employees;*
- Check warehouse and metastore

---

# Hive - Partitions - Summary

---

- To avoid the full table scan
- The data is stored in different files in warehouse defined by the partitions
- Define the partitions using “partition by” in “create table”
- We can also add a partition later
- Partition can happen on multiple columns (year=2012, month=10, day=12)



---

# Hive - Views

---

- *SELECT \* FROM employees where department='Engineering';*
- Create a view

*CREATE VIEW employees\_engineering AS  
SELECT \* FROM employees where department='Engineering';*

- Now query from the view

*SELECT \* FROM employees\_engineering;*

---

# Hive - Views - Summary

---

- Allows a query to be saved and treated like a table
- Logical construct - does not store data
- Hides the query complexity
- Divide long and complicated query into smaller and manageable pieces
- Similar to writing a function in a programming language

---

# Hive - Load JSON Data

---

- Download [JSON-SERDE BINARIES](#)
- ADD JAR

*hdfs:///data/serde/json-serde-1.3.6-SNAPSHOT-jar-with-dependencies.jar;*

- Create Table

*CREATE EXTERNAL TABLE tweets\_raw (*

*)*

***ROW FORMAT SERDE 'org.apache.hive.hcatalog.data.JsonSerDe'***

***LOCATION '/user/abhinav9884/senti/upload/data/tweets\_raw';***

---

# Hive - Sorting & Distributing - Order By

---

## **ORDER BY x**

- Guarantees global ordering
- Data goes through just one reducer
- This is unacceptable for large datasets as it will overload the reducer
- You end up one sorted file as output



---

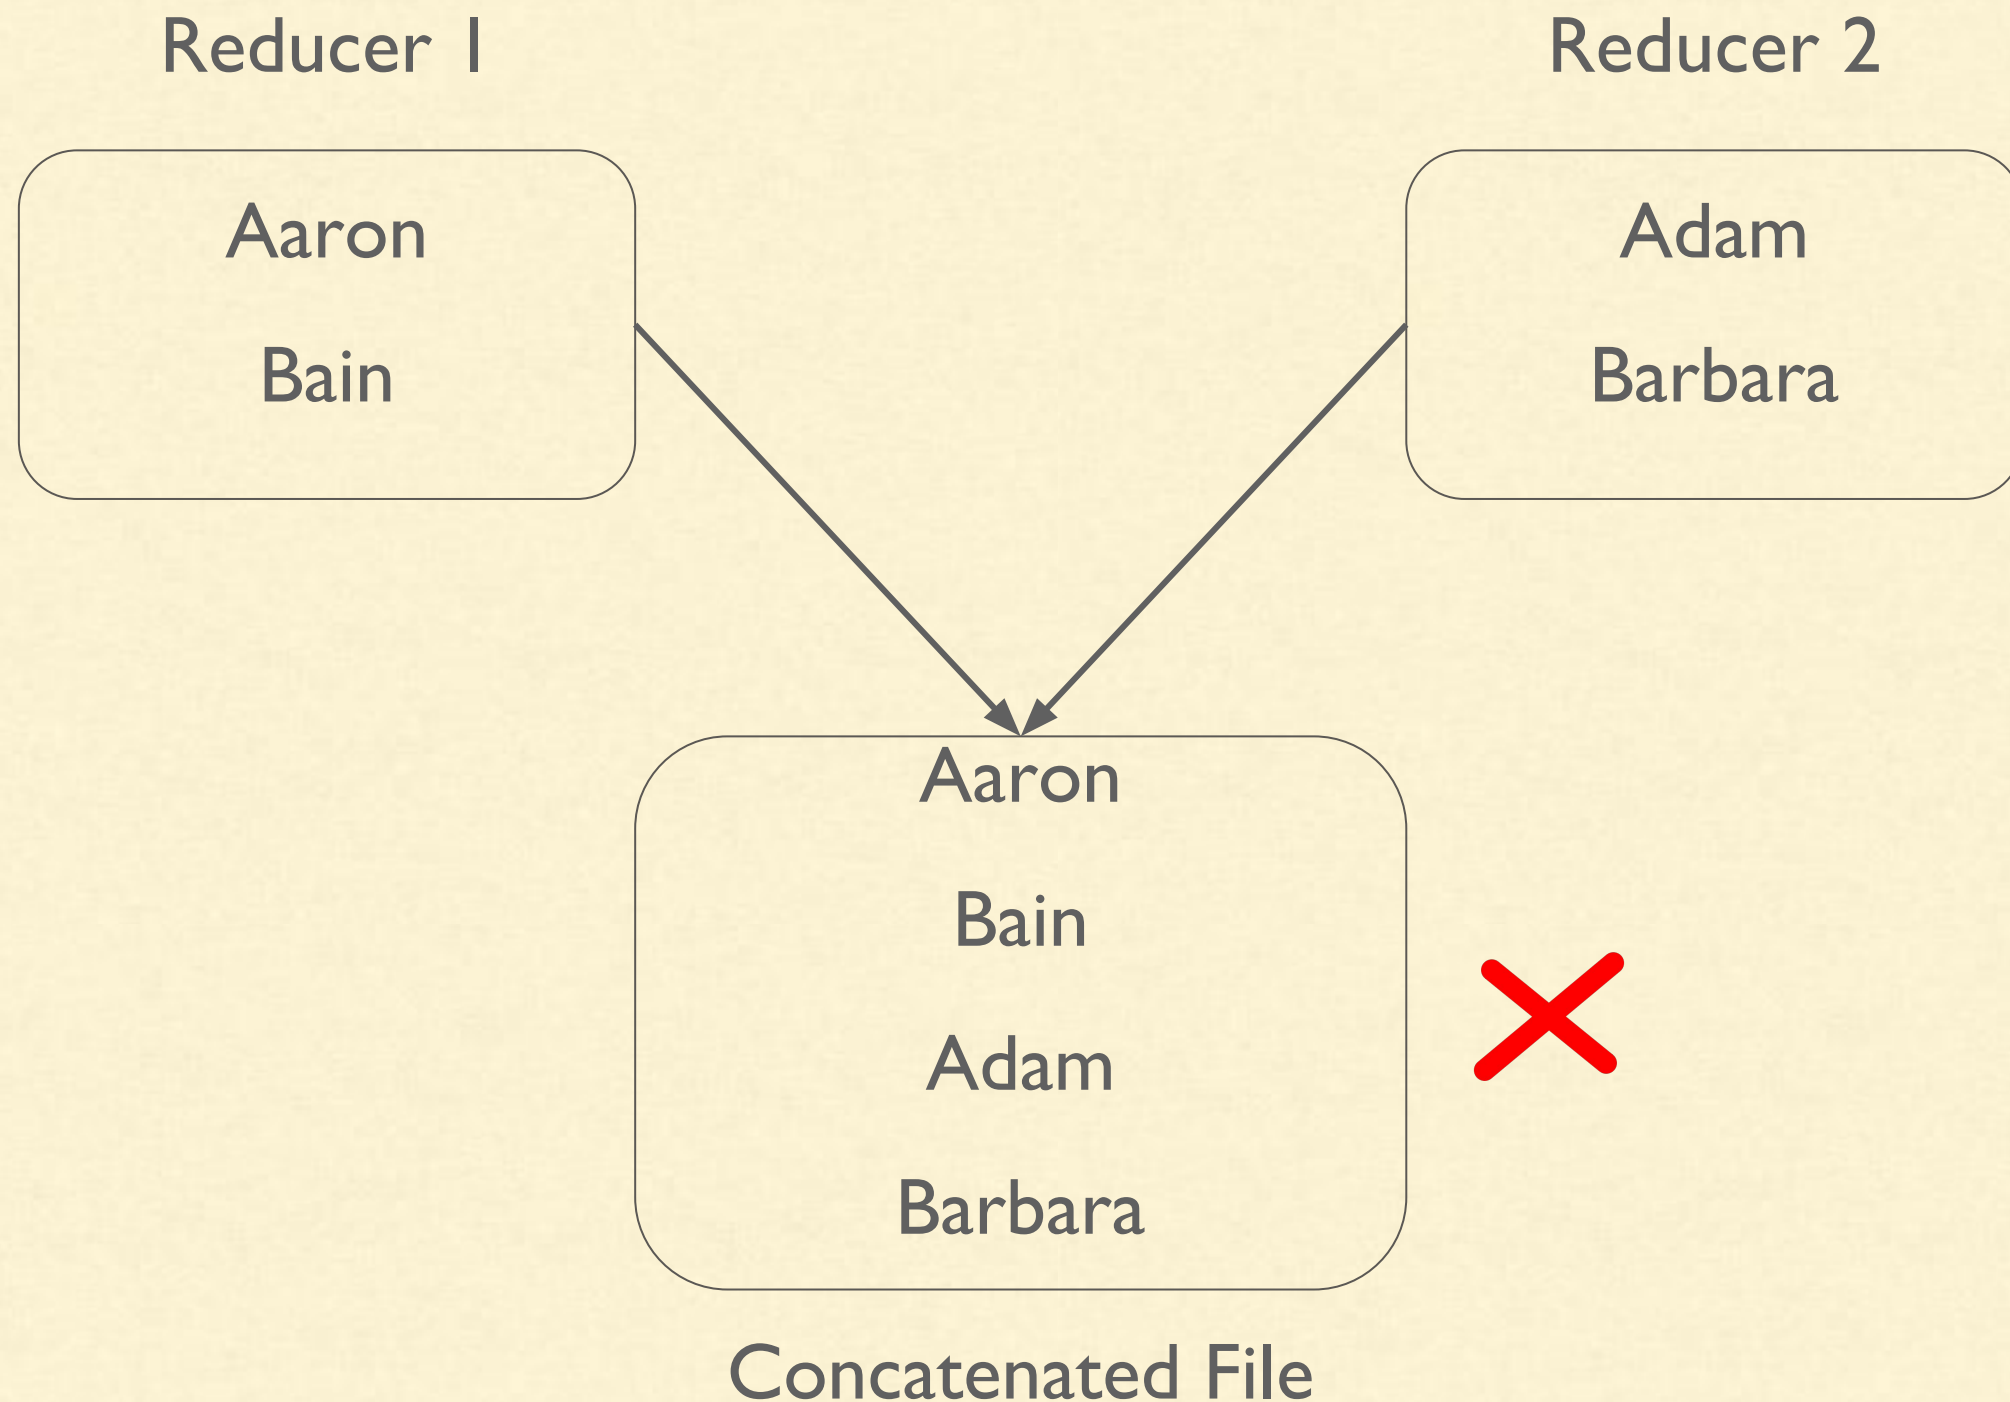
# Hive - Sorting & Distributing - Sort By

---

## **SORT BY x**

- Orders data at each of N reducers
- Number of reducers are 1 per 1GB
- You end up with N or more sorted files with overlapping ranges

# Hive - Sorting & Distributing - Sort By



---

# Hive - Sorting & Distributing - Distribute By

---

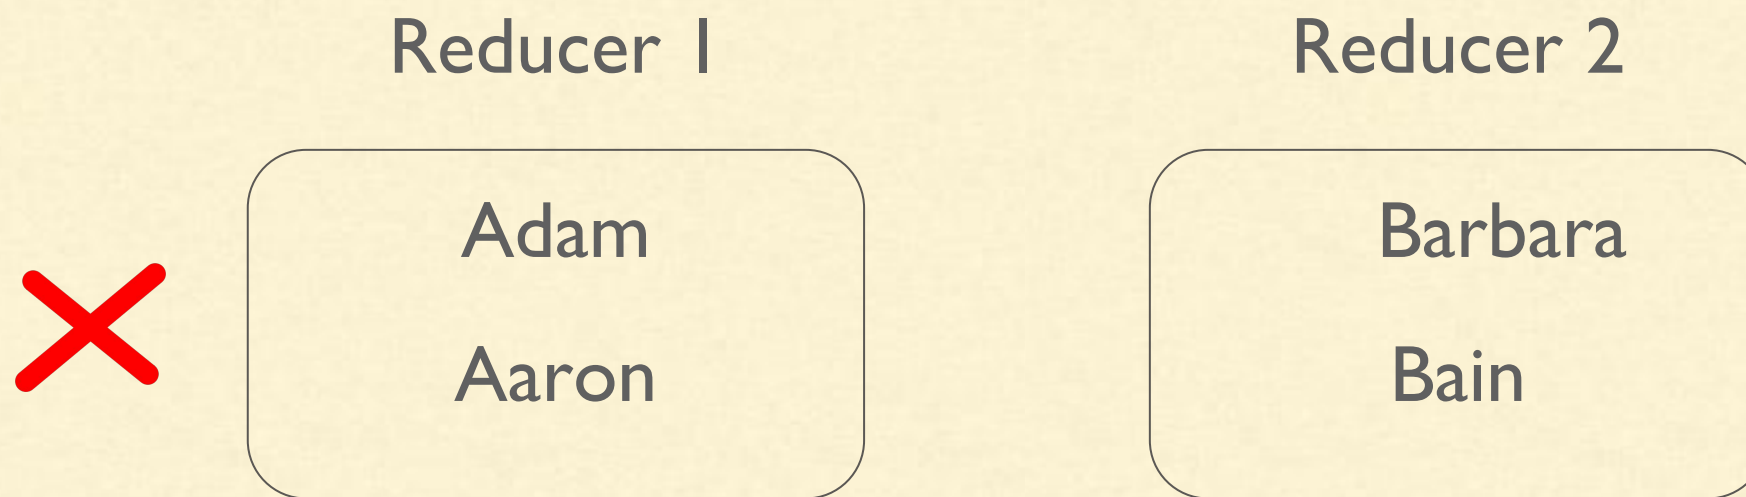
## **DISTRIBUTE BY x**

- Ensures each of N reducers gets non-overlapping ranges of x
- But doesn't sort the output of each reducer
- You end up with N or unsorted files with non-overlapping ranges

---

# Hive - Sorting & Distributing - Distribute By

---





---

# Hive - Sorting & Distributing - Cluster By

---

## **CLUSTER BY x**

- Gives global ordering
- Is the same as (DISTRIBUTE BY x and SORT BY x)
- CLUSTER BY is basically the more scalable version of ORDER BY

---

# Hive - Bucketing

---

```
CREATE TABLE page_view(viewTime INT, userid BIGINT,  
    page_url STRING, referrer_url STRING,  
    ip STRING COMMENT 'IP Address of the User'  
)  
COMMENT 'This is the page view table'  
PARTITIONED BY(dt STRING, country STRING)  
CLUSTERED BY(userid) INTO 32 BUCKETS  
ROW FORMAT DELIMITED FIELDS TERMINATED BY '\001'  
COLLECTION ITEMS TERMINATED BY '\002'  
MAP KEYS TERMINATED BY '\003'  
STORED AS SEQUENCEFILE;
```

---

# Hive - ORC Files

---

- Optimized Row Columnar file format
- Provides a highly efficient way to store Hive data
- Improves performance when
  - Reading
  - Writing
  - Processing
- Has a built-in index, min/max values, and other aggregations
- Proven in large-scale deployments
  - Facebook uses the ORC file format for a 300+ PB deployment

---

# Hive - ORC Files - Example

---

```
CREATE TABLE orc_table (  
    first_name STRING,  
    last_name STRING  
) STORED AS ORC;
```

```
INSERT INTO orc_table VALUES('John', 'Gill');  
SELECT * from orc_table;
```

To Know more, please visit <https://orc.apache.org>



---

# Hive - Quick Recap

---

- Each table has got a location
- By default the table is in a directory under the location  
/apps/hive/warehouse
- We can override that location by mentioning 'location' in create table clause
- Load data copies the data if it is local
- Load moves the data if it is on hdfs for both external and managed tables
- Dropping managed table deletes the data at the 'location'
- Dropping external table does not delete the data at the 'location'
- The metadata is stored in the relational database - hive metastore

---

# Hive - Connecting to Tableau

---

- Tableau is a visualization tool
- Tableau allows for instantaneous insight by transforming data into visually appealing, interactive visualizations called dashboards

---

# Hive - Connecting to Tableau - Steps

---

- Download and install Tableau desktop from <https://www.tableau.com/products/desktop>
- Download and install Hortonworks ODBC driver for Apache Hive for your OS  
<https://hortonworks.com/downloads/>

---

# Hive - Connecting to Tableau - Hands-on

---

Visualize top 10 stocks with highest opening price on Dec 31, 2009



---

# Hive - Quick Demo

---

1. **Copy data from /data/ml100k/u.data into our hdfs home**
2. **Open Hive in Hue and run following:**

```
CREATE TABLE u_data( userid INT, movieid INT, rating INT, unixtime STRING)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY '\t'  
STORED AS TEXTFILE;
```

```
LOAD DATA INPATH '/user/sandeepgiri9034/u.data' overwrite into table u_data;
```

```
select * from u_data limit 5;
```

```
select movieid, avg(rating) ar from u_data group by movieid order by ar desc
```

---

# Hive - Quick Demo

---

## *Join with Movie Names*

*create view top100m as  
select movieid, avg(rating) ar from u\_data group by movieid order by ar desc*

*CREATE TABLE m\_data( movieid INT, name STRING)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY '|'  
STORED AS TEXTFILE;*

*load data inpath '/user/sandeepgiri9034/u.item' into table m\_data;  
select \* from m\_data limit 100  
select \* from m\_data, top100m where top100m.movieid = m\_data.movieid*

---

# Hive - Assignment

---

1. For each movie how many users rated it
2. For movies having more than 30 ratings, what is the average rating