



Loading + Saving Data







Loading & Saving Data

1. So far we
 - a. either converted in-memory data
 - b. Or used the HDFS file



Loading & Saving Data

1. So far we
 - a. either converted in-memory data
 - b. Or used the HDFS file
2. Spark supports wide variety of datasets
3. Can access data through InputFormat & OutputFormat
 - a. The interfaces used by Hadoop
 - b. Available for many common file formats and storage systems (e.g., S3, HDFS, Cassandra, HBase, etc.).



Common Data Sources

File formats  	Stores  

Common Data Sources

File formats 	Stores 
<ul style="list-style-type: none">• Text, JSON, SequenceFiles, Protocol buffers.• We can also configure compression	

Common Data Sources

File formats 	Stores 
<ul style="list-style-type: none">• Text, JSON, SequenceFiles, Protocol buffers.• We can also configure compression	<p>Filesystems</p> <ul style="list-style-type: none">• Local, NFS, HDFS, Amazon S3 <p>Databases and key/value stores</p> <ul style="list-style-type: none">• For Cassandra, HBase, Elasticsearch, and JDBC databases.

Loading & Saving Data



Structured data sources through Spark SQL aka Data Frames

- + Efficient API for structured data sources, including JSON and Apache Hive
- + Covered later

Common supported file formats

- + A file could be in any format
- + If we know upfront, we can read it and load it
- + Else we can use “file” command like tool

```
[sandeep@ip-172-31-60-179 ~]$ file spark-2.0.2-bin-hadoop2.7.tgz
spark-2.0.2-bin-hadoop2.7.tgz: gzip compressed data, from Unix, last modified: Tue Nov  8 01:58:16 2016
[sandeep@ip-172-31-60-179 ~]$
```


Common supported file formats

Text files

- + Very common. Plain old text files. Printable chars
- + Records are assumed to be one per line.
- + Unstructured Data

Example:

```
[sandeep@ip-172-31-60-179 ~]$ head /cxldata/big.txt
The Project Gutenberg EBook of The Adventures of Sherlock Holmes
by Sir Arthur Conan Doyle
(#15 in our series by Sir Arthur Conan Doyle)

Copyright laws are changing all over the world. Be sure to check the
copyright laws for your country before downloading or redistributing
this or any other Project Gutenberg eBook.

This header should be the first thing seen when viewing this Project
Gutenberg file. Please do not remove it. Do not change or edit the
```

Common supported file formats

JSON files

- + Javascript Object Notation
- + Common text-based format
- + Semistructured; most libraries require one record per line.

Example:

```
{  
  "name" : "John",  
  "age" : 31,  
  "knows" : ["C", "C++"]  
}
```

Common supported file formats

CSV files

- + Very common text-based format
- + Often used with spreadsheet applications.
- + Comma separated Values

Example:

```
Title, Author, ISBN13, Pages
1984, George Orwell, 978-0451524935, 268
Animal Farm, George Orwell, 978-0451526342, 144
Brave New World, Aldous Huxley, 978-0060929879, 288
Fahrenheit 451, Ray Bradbury, 978-0345342966, 208
Jane Eyre, Charlotte Brontë, 978-0142437209, 532
Wuthering Heights, Emily Brontë, 978-0141439556, 416
```

Common supported file formats

Sequence files

- + Compact Hadoop file format used for key/value data.
- + Key and values can be binary data
- + To bundle together many small files

Record Structure



Block Structure



See More at <https://wiki.apache.org/hadoop/SequenceFile>

Common supported file formats

Protocol buffers

- + A fast, space-efficient multilanguage format.
- + More compact than JSON.

```
message Person {  
    required string name = 1;  
    required int32 id = 2;  
    optional string email = 3;  
}
```

See More at <https://developers.google.com/protocol-buffers/>

Common supported file formats

Object Files

- + For data from a Spark job to be consumed by another
- + Breaks if you change your classes - Java Serialization.

Handling Text Files - scala



Loading Files

```
var input = sc.textFile("/data/ml-100k/u1.test")
```



Handling Text Files - scala

Loading Files

```
var input = sc.textFile("/data/ml-100k/u1.test")
```

Loading Directories

```
var input = sc.wholeTextFiles("/data/ml-100k");  
var lengths = input.mapValues(x => x.length);  
lengths.collect();
```

```
[(u'hdfs://ip-172-31-53-48.ec2.internal:8020/data/ml-100k/mku.sh', 643),  
(u'hdfs://ip-172-31-53-48.ec2.internal:8020/data/ml-100k/u.data', 1979173),  
(u'hdfs://ip-172-31-53-48.ec2.internal:8020/data/ml-100k/u.genre', 202),  
(u'hdfs://ip-172-31-53-48.ec2.internal:8020/data/ml-100k/u.info', 36) ...]
```




Handling Text Files - scala

Loading Files

```
var input = sc.textFile("/data/ml-100k/u1.test")
```

Loading Directories

```
var input = sc.wholeTextFiles("/data/ml-100k");  
var lengths = input.mapValues(x => x.length);  
lengths.collect();
```

```
[(u'hdfs://ip-172-31-53-48.ec2.internal:8020/data/ml-100k/mku.sh', 643),  
(u'hdfs://ip-172-31-53-48.ec2.internal:8020/data/ml-100k/u.data', 1979173),  
(u'hdfs://ip-172-31-53-48.ec2.internal:8020/data/ml-100k/u.genre', 202),  
(u'hdfs://ip-172-31-53-48.ec2.internal:8020/data/ml-100k/u.info', 36) ...]
```

Saving Files

```
lengths.saveAsTextFile(outputDir)
```

Comma / Tab -Separated Values (CSV / TSV)

1. Records are stored one per line,
2. Fixed number of fields per line
3. Fields are separated by a comma (tab in TSV)
4. We get row number to detect header etc.

Loading CSV - Sample Data



Data: </data/spark/temps.csv>

20, NYC, 2014-01-01
20, NYC, 2015-01-01
21, NYC, 2014-01-02
23, BLR, 2012-01-01
25, SEATTLE, 2016-01-01
21, CHICAGO, 2013-01-05
24, NYC, 2016-5-05



Loading CSV - Simple Approach

```
var lines = sc.textFile("/data/spark/temps.csv");  
var recordsRDD = lines.map(line => line.split(","));  
recordsRDD.take(10);
```

```
Array(  
  Array(20, " NYC", " 2014-01-01"),  
  Array(20, " NYC", " 2015-01-01"),  
  Array(21, " NYC", " 2014-01-02"),  
  Array(23, " BLR", " 2012-01-01"),  
  Array(25, " SEATTLE", " 2016-01-01"),  
  Array(21, " CHICAGO", " 2013-01-05"),  
  Array(24, " NYC", " 2016-5-05")  
)
```


Loading CSV - Example



spark-shell --packages net.sf.opencsv:opencsv:2.3

Or

Add this to sbt: *libraryDependencies += "net.sf.opencsv" % "opencsv" % "2.3"*

```
import au.com.bytecode.opencsv.CSVParser

var a = sc.textFile("/data/spark/temps.csv");
var p = a.map(
  line => {
    val parser = new CSVParser(',')
    parser.parseLine(line)
  })
p.take(1)
//Array(Array(20, " NYC", " 2014-01-01"))
```

<https://gist.github.com/girisandeep/b721cf93981c338665c328441d419253>

Loading CSV - Example Efficient



<https://gist.github.com/girisandeep/fddf49ef97fde429a0d3256160b257c1>



Loading CSV - Example Efficient

```
import au.com.bytecode.opencsv.CSVParser  
var linesRdd = sc.textFile("/data/spark/temps.csv");
```

<https://gist.github.com/girisandeep/fddf49ef97fde429a0d3256160b257c1>

Loading CSV - Example Efficient



```
import au.com.bytecode.opencsv.CSVParser
var linesRdd = sc.textFile("/data/spark/temps.csv");
def parseCSV(itr:Iterator[String]):Iterator[Array[String]] = {
  val parser = new CSVParser(',')
  for(line <- itr)
    yield parser.parseLine(line)
}
```

<https://gist.github.com/girisandeep/fddf49ef97fde429a0d3256160b257c1>

Loading CSV - Example Efficient



```
import au.com.bytecode.opencsv.CSVParser
var linesRdd = sc.textFile("/data/spark/temps.csv");
def parseCSV(itr:Iterator[String]):Iterator[Array[String]] = {
  val parser = new CSVParser(',')
  for(line <- itr)
    yield parser.parseLine(line)
}
//Check with simple example
val x = parseCSV(Array("1,2,3","a,b,c").iterator)
val result = linesRdd.mapPartitions(parseCSV)
```

<https://gist.github.com/girisandeep/fddf49ef97fde429a0d3256160b257c1>



Loading CSV - Example Efficient

```
import au.com.bytecode.opencsv.CSVParser
var linesRdd = sc.textFile("/data/spark/temps.csv");
def parseCSV(itr:Iterator[String]):Iterator[Array[String]] = {
  val parser = new CSVParser(',')
  for(line <- itr)
    yield parser.parseLine(line)
}
//Check with simple example
val x = parseCSV(Array("1,2,3","a,b,c").iterator)
val result = linesRdd.mapPartitions(parseCSV)
result.take(1)
//Array[Array[String]] = Array(Array(20, " NYC", " 2014-01-01"))
```

<https://gist.github.com/girisandeep/fddf49ef97fde429a0d3256160b257c1>

Tab Separated Files

Similar to csv:

```
val parser = new CSVParser('\t')
```

SequenceFiles

- Popular Hadoop format
 - For handling small files
 - Create InputSplits without too much transport

SequenceFiles

- Popular Hadoop format
 - For handling small files
 - Create InputSplits without too much transport
- Composed of flat files with key/value pairs.
- Has Sync markers
 - Allow to seek to a point
 - Then resynchronize with the record boundaries
 - Allows Spark to efficiently read in parallel from multiple nodes

Loading SequenceFiles

```
val data = sc.sequenceFile(inFile,  
"org.apache.hadoop.io.Text", "org.apache.hadoop.io.IntWritable")  
data.map(func)  
...
```

```
data.saveAsSequenceFile(outputFile)
```

Saving SequenceFiles - Example

```
var rdd = sc.parallelize(Array(("key1", 1.0), ("key2", 2.0), ("key3", 3.0)))  
rdd.saveAsSequenceFile("pysequencefile1")
```

Saving SequenceFiles - Example

```
var rdd = sc.parallelize(Array(("key1", 1.0), ("key2", 2.0), ("key3", 3.0)))  
rdd.saveAsSequenceFile("pysequencefile1")
```

```
[sandeep@ip-172-31-60-179 ~]$ hadoop fs -ls pysequencefile1/
```

```
Found 5 items
```

-rw-r--r--	3	sandeep	hdfs	0	2017-06-21	21:52	pysequencefile1/_SUCCESS
-rw-r--r--	3	sandeep	hdfs	88	2017-06-21	21:52	pysequencefile1/part-00000
-rw-r--r--	3	sandeep	hdfs	109	2017-06-21	21:52	pysequencefile1/part-00001
-rw-r--r--	3	sandeep	hdfs	109	2017-06-21	21:52	pysequencefile1/part-00002
-rw-r--r--	3	sandeep	hdfs	109	2017-06-21	21:52	pysequencefile1/part-00003

Loading SequenceFiles - Example

```
import org.apache.hadoop.io.DoubleWritable  
import org.apache.hadoop.io.Text
```

Loading SequenceFiles - Example

```
import org.apache.hadoop.io.DoubleWritable
import org.apache.hadoop.io.Text

val myrdd = sc.sequenceFile(
  "pysequencefile1",
  classOf[Text], classOf[DoubleWritable])
```

Loading SequenceFiles - Example

```
import org.apache.hadoop.io.DoubleWritable
import org.apache.hadoop.io.Text

val myrdd = sc.sequenceFile(
  "pysequencefile1",
  classOf[Text], classOf[DoubleWritable])

val result = myrdd.map{case (x, y) => (x.toString, y.get())}
result.collect()

//Array((key1,1.0), (key2,2.0), (key3,3.0))
```

Object Files

- Simple wrapper around SequenceFiles
- Values are written out using Java Serialization.
- Intended to be used for Spark jobs communicating with other Spark jobs
- Can also be quite slow.

Object Files

- Saving - `saveAsObjectFile()` on an RDD
- Loading - `objectFile()` on `SparkContext`
- Require almost no work to save almost arbitrary objects.
- Not available in python using pickle file instead
- If you change the objects, old files may not be valid

Pickle File

- Python way of handling object files
- Uses Python's pickle serialization library
- Saving - `saveAsPickleFile()` on an RDD
- Loading - `pickleFile()` on SparkContext
- Can also be quite slow as Object Files

Non-filesystem data sources - hadoopFile

- Access Hadoop-supported storage formats
- Many key/value stores provide Hadoop input formats
- Example providers: HBase, MongoDB

- Older: `hadoopFile()` / `saveAsHadoopFile()`
- Newer: `newAPIHadoopDataset()` / `saveAsNewAPIHadoopDataset()`
- Takes a `Configuration` object on which you set the Hadoop properties

Hadoop Input and Output Formats - Old API

```
hadoopFile(path, inputFormatClass, keyClass, valueClass, keyConverter=None,  
           valueConverter=None, conf=None, batchSize=0)
```

Read an 'old' Hadoop InputFormat with arbitrary key and value class from HDFS, a local file system (available on all nodes), or any Hadoop-supported file system URI. The mechanism is the same as for `sc.sequenceFile`.

A Hadoop configuration can be passed in as a Python dict. This will be converted into a Configuration in Java.

Parameters:

path – path to Hadoop file

inputFormatClass – fully qualified classname of Hadoop InputFormat (e.g. “org.apache.hadoop.mapred.TextInputFormat”)

keyClass – fully qualified classname of key Writable class (e.g. “org.apache.hadoop.io.Text”)

valueClass – fully qualified classname of value Writable class (e.g. “org.apache.hadoop.io.LongWritable”)

keyConverter – (None by default)

valueConverter – (None by default)

conf – Hadoop configuration, passed in as a dict (None by default)

batchSize – The number of Python objects represented as a single Java object. (default 0, choose batchSize automatically)

Hadoop Input and Output Formats - New API

```
newAPIHadoopFile(path, inputFormatClass, keyClass, valueClass, keyConverter=None,
                  valueConverter=None, conf=None, batchSize=0)
```

Read a 'new API' Hadoop InputFormat with arbitrary key and value class from HDFS, a local file system (available on all nodes), or any Hadoop-supported file system URI. The mechanism is the same as for `sc.sequenceFile`.

A Hadoop configuration can be passed in as a Python dict. This will be converted into a Configuration in Java

Parameters:

path – path to Hadoop file

inputFormatClass – fully qualified classname of Hadoop InputFormat (e.g.

“org.apache.hadoop.mapreduce.lib.input.TextInputFormat”)

keyClass – fully qualified classname of key Writable class (e.g. “org.apache.hadoop.io.Text”)

valueClass – fully qualified classname of value Writable class (e.g. “org.apache.hadoop.io.LongWritable”)

keyConverter – (None by default)

valueConverter – (None by default)

conf – Hadoop configuration, passed in as a dict (None by default)

batchSize – The number of Python objects represented as a single Java object. (default 0, choose batchSize automatically)

Hadoop Input and Output Formats - Old API

Loading Data from mongodb

```
See https://databricks.com/blog/2015/03/20/using-mongodb-with-spark.html  
# set up parameters for reading from MongoDB via Hadoop input format  
config = {"mongo.input.uri": "mongodb://localhost:27017/marketdata.minbars"}  
inputFormatClassName = "com.mongodb.hadoop.MongoInputFormat"  
  
keyClassName = "org.apache.hadoop.io.Text"  
valueClassName = "org.apache.hadoop.io.MapWritable"  
  
# read the 1-minute bars from MongoDB into Spark RDD format  
minBarRawRDD = sc.newAPIHadoopRDD(inputFormatClassName, keyClassName,  
valueClassName, None, None, config)
```

Protocol buffers

- Developed at Google for internal RPCs
- Open sourced
- Structured data - fields & types of fields defined
- Fast for encoding and decoding (20-100x than XML)
- Take up the minimum space (3-10x than xml)
- Defined using a domain-specific language
- Compiler generates accessor methods in variety of languages
- Consist of fields: optional, required, or repeated
- While parsing
 - A missing optional field => success
 - A missing required field => failure
- So, make new fields as optional (remember object file failures?)

Protocol buffers - Example

```
package tutorial;
message Person {
  required string name = 1;
  required int32 id = 2;
  optional string email = 3;
  enum PhoneType {
    MOBILE = 0;
    HOME = 1;
    WORK = 2;
  }
  message PhoneNumber {
    required string number = 1;
    optional PhoneType type = 2 [default = HOME];
  }
  repeated PhoneNumber phone = 4;
}

message AddressBook {
  repeated Person person = 1;
}
```


Protocol buffers - Steps

1. [Download](#) and install protocol buffer compiler
2. `pip install protobuf`
3. `protoc -I=$SRC_DIR --python_out=$DST_DIR $SRC_DIR/addressbook.proto`
4. create objects
5. Convert those into protocol buffers
6. See [this project](#)

File Compression

1. To Save Storage & Network Overhead
2. With most hadoop output formats we can specify compression codecs
3. Compression should not require the whole file at once
4. Each worker can find start of record => splittable
5. You can configure HDP for LZO using Ambari:
http://docs.hortonworks.com/HDPDocuments/Ambari-2.2.2.0/bk_ambari_reference_guide/content/_configure_core-sitexml_for_lzo.html

File Compression Options

Characteristics of compression:

- Splittable
- Speed
- Effectiveness on Text
- Code

File Compression Options

Format	Splittable	Speed	Effectiveness on text	Hadoop compression codec	comments
gzip	N	Fast	High	org.apache.hadoop.io.com.GzipCodec	
<i>lzo</i>	Y	<i>V. Fast</i>	<i>Medium</i>	<i>com.hadoop.compression.lzo.LzoCodec</i>	<i>LZO requires installation on every worker node</i>
<i>bzip2</i>	Y	<i>Slow</i>	<i>V. High</i>	<i>org.apache.hadoop.io.com.BZip2Codec</i>	<i>Uses pure Java for splittable version</i>
zlib	N	Slow	Medium	org.apache.hadoop.io.com.DefaultCodec	Default compression codec for Hadoop
Snappy	N	V. Fast	Low	org.apache.hadoop.io.com.SnappyCodec	There is a pure Java port of Snappy but it is not yet available in Spark/Hadoop

Handling LZO

1. Enable in HADOOP by updating the conf of hadoop
http://docs.hortonworks.com/HDPDocuments/Ambari-2.2.2.0/bk_ambari_reference_guide/content/_configure_core-sitexml_for_lzo.html

2. Create data:
`$ bzip2 --stdout file.bz2 | lzop -o file.lzo`

3. Update Spark-env.sh with
`export SPARK_CLASSPATH=$SPARK_CLASSPATH:hadoop-lzo-0.4.20-SNAPSHOT.jar`

In your code, use:
`conf.set("io.compression.codecs", "com.hadoop.compression.lzo.LzopCodec");`

Ref: <https://gist.github.com/zedar/c43cbc7ff7f98abee885>

Loading + Saving Data: File Systems

Local/“Regular” FS

1. `rdd = sc.textFile("file:///home/holden/happypandas.gz")`
2. The path has to be available on all nodes.
Otherwise, load it locally and distribute using `sc.parallelize`

Amazon S3

1. Popular option
2. Good if nodes are inside EC2
3. Use path in all input methods (textFile, hadoopFile etc)
s3n://bucket/path-within-bucket
4. Set Env. Vars: `AWS_ACCESS_KEY_ID` `AWS_SECRET_ACCESS_KEY`

More details: <https://cloudxlab.com/blog/access-s3-files-spark/>

HDFS

1. The Hadoop Distributed File System
2. Spark and HDFS can be colocated on the same machines
3. Spark can take advantage of this data locality to avoid network overhead
4. In all i/o methods, use path: `hdfs://master:port/path`
5. Use only the version of spark w.r.t HDFS version



Spark - Loading + Saving Data

Thank you!

