



SparkR



SparkR (R on Spark)

“SparkR is an R package that provides light-weight frontend to use Apache Spark on R”

- Distributed data frame - supports
 - selection, filtering, aggregation etc
- Can Handle large datasets
- Supports distributed machine learning using MLlib

SparkR DataFrames

- A DataFrame is a distributed collection of data organized into named columns
- Equivalent to a table in a relational database or a data frame in R, but with richer optimizations under the hood
- Can be constructed from a wide array of sources such as: structured data files, tables in Hive, external databases, or existing local R data frames

Launch SparkR

```
# Login to CloudeXLab web console  
/usr/spark2.0.1/bin/sparkR
```

Creating DataFrames - From local dataframes

faithful - R Dataframe - waiting time between eruptions and the duration of the eruption

Creating DataFrames - From local dataframes

faithful - R Dataframe - waiting time between eruptions and the duration of the eruption

```
df = createDataFrame(spark, faithful)
```

Creating DataFrames - From local dataframes

faithful - R Dataframe - waiting time between eruptions and the duration of the eruption

```
df = createDataFrame(spark, faithful)
```

```
# Displays the content of the DataFrame to stdout  
head(df)
```

##	eruptions	waiting
##1	3.600	79
##2	1.800	54
##3	3.333	74

Data Frame Operations

Selecting Rows and Columns

```
# Select only the "eruptions" column
```

```
> res = select(df, df$eruptions)
```

```
> head(res)
```

```
eruptions
```

```
1      3.600
```

```
2      1.800
```

```
3      3.333
```

```
4      2.283
```

```
5      4.533
```

```
6      2.883
```

```
> # You can also pass in column name as strings
```

```
> head(select(df, "eruptions"))
```


Data Frame Operations

Selecting Rows and Columns

```
# Filter the DataFrame to only  
# retain rows with wait times shorter than 50 mins  
  
> res = filter(df, df$waiting < 50)  
> head(res)  
  eruptions waiting  
1    1.7500      47  
2    1.7500      47  
3    1.8667      48  
4    1.7500      48  
5    2.1667      48  
6    2.1000      49
```

Data Frame Operations

Grouping and Aggregation

```
# We use the `n` operator to count the number of times  
# each waiting time appears
```

```
> grpd = groupBy(df, df$waiting)  
> N = n(df$waiting)  
> res = summarize(grpd, count = N)  
> head(res)
```

	waiting	count
1	70	4
2	67	1
3	69	2
4	88	6
5	49	5
6	64	4

Data Frame Operations

Grouping and Aggregation

```
# We use the `n` operator to count the number of times  
# each waiting time appears
```

```
> head(summarize(groupBy(df, df$waiting), count =  
n(df$waiting)))
```

	waiting	count
1	70	4
2	67	1
3	69	2
4	88	6
5	49	5
6	64	4

Data Frame Operations

Sorting

We can also sort the output from the aggregation to get the most common waiting times

```
> waiting_counts = summarize(groupBy(df, df$waiting),  
count = n(df$waiting))  
> head(arrange(waiting_counts,  
desc(waiting_counts$count)))
```

	waiting	count
1	78	15
2	83	14
3	81	13
4	77	12
5	82	12
6	79	10

Data Frame Operations

Operating on Columns

```
# Convert waiting time from hours to seconds.  
# Note that we can assign this to a new column in the same DataFrame
```

```
df$waiting_secs = df$waiting * 60  
head(df)
```

	eruptions	waiting	waiting_secs
1	3.600	79	4740
2	1.800	54	3240
3	3.333	74	4440
4	2.283	62	3720
5	4.533	85	5100
6	2.883	55	3300

Creating DataFrames - From JSON

```
$ hadoop fs -cat /data/spark/people.json
```

```
{"name": "Michael"}  
{"name": "Andy", "age": 30}  
{"name": "Justin", "age": 19}
```

Creating DataFrames - From JSON

```
$ hadoop fs -cat /data/spark/people.json
```

```
{ "name": "Michael" }  
{ "name": "Andy", "age": 30 }  
{ "name": "Justin", "age": 19 }
```

```
$ /usr/spark2.0.1/bin/sparkR
```

```
> people = read.df(spark, "/data/spark/people.json", "json")
```

Creating DataFrames - From JSON

```
$ hadoop fs -cat /data/spark/people.json
```

```
{"name": "Michael"}  
{"name": "Andy", "age": 30}  
{"name": "Justin", "age": 19}
```

```
$ /usr/spark2.0.1/bin/sparkR
```

```
> people = read.df(spark, "/data/spark/people.json", "json")  
> head(people)
```

	age	name
1	NA	Michael
2	30	Andy
3	19	Justin

Running SQL Queries from SparkR

```
# Load a JSON file
people = read.df(spark, "/data/spark/people.json", "json")

# Register this DataFrame as a table.
createOrReplaceTempView(people, "peopleview")

# SQL statements can be run by using the sql method

teenagers = sql(spark, "SELECT name FROM peopleview WHERE age >= 13 AND
age <= 19")

head(teenagers)

      name
1 Justin
```



SparkR

Thank you!



Creating DataFrames from JSON

Example:

In Scala:

```
var df = spark.read.json("/data/spark/people.json")
```

```
# Displays the content of the DataFrame to stdout  
df.show()
```

Or In R:

```
df <- read.json("/data/spark/people.json")  
showDF(df)
```

```
{"name":"Michael"}  
{"name":"Andy", "age":30}  
{"name":"Justin", "age":19}
```

Data Sources

- Spark SQL supports operating on a variety of data sources through the DataFrame interface.
- A DataFrame can be operated on as normal RDDs and can also be registered as a temporary table.
- Registering a DataFrame as a table allows you to run SQL queries over its data.

Beeline

1. `/usr/spark2.0.1/bin/beeline`
2. `!connect jdbc:hive2://c.cloudxlab.com:10000`
3. `use sg;`
4. `show tables;`
5. `select * from employees;`

Creating DataFrames - From JSON

```
$ hadoop fs -cat /data/spark/people.json
```

```
{ "name": "Michael" }  
{ "name": "Andy", "age": 30 }  
{ "name": "Justin", "age": 19 }
```

```
$ /usr/spark2.0.1/bin/sparkR
```

```
> people = read.df(spark, "/data/spark/people.json", "json")
```

```
# SparkR automatically infers the schema from the JSON file
```

```
> printSchema(people)
```

```
# root
```

```
# |-- age: integer (nullable = true)
```

```
# |-- name: string (nullable = true)
```