

In [1]: `cd G:/`

G:\

In [2]: `cd smart_bridge`

G:\smart_bridge

In [3]: `import pandas as pd
import matplotlib.pyplot as plt
import numpy as np`

In [4]: `df = pd.read_csv("SENTIMENTAL ANALYSIS BASED ON FLIGHT REVIEWS.csv")
df.head()`

Out[4]:

	airline_sentiment	airline	text
0	positive	Virgin America	@VirginAmerica What @dhepburn said.
1	positive	Virgin America	@VirginAmerica plus you've added commercials t...
2	positive	Virgin America	@VirginAmerica I didn't today... Must mean I n...
3	negative	Virgin America	@VirginAmerica it's really aggressive to blast...
4	negative	Virgin America	@VirginAmerica and it's a really big bad thing...

In [5]: `df.airline_sentiment = df.airline_sentiment.map({"positive":1, "negative":0})
df.airline = df.airline.map({"Virgin America":2, "United":1})`

In [6]: `from nltk.stem.snowball import SnowballStemmer
from nltk.stem import PorterStemmer
from nltk.corpus import stopwords
import re
import nltk
nltk.download('stopwords')
nltk.download('wordnet')`

[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\Admin\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\Admin\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!

Out[6]: True

```
In [7]: my_stopwords = ["is", "am", "an", "are", "of", "can", "were", "he", "she", "him",
                        "also", "to", "for", "the", "we", "you", "by", "during", "on", "in",
                        "that", "would", "this", "there", "youd", "who", "then", "they", "it", "a", "an"]

def clean_review(review):
    review=re.sub("(n't|(\w+nt))", " not", review)
    review=re.sub("(\\@\\w+)", "", review)
    review=re.sub("[^a-zA-Z]", " ", review)
    review=re.sub("\\s\\s", " ", review)

    review=review.lower()
    review=review.split()
    ps=PorterStemmer()
    lemma = nltk.wordnet.WordNetLemmatizer()

    review=[ps.stem(lemma.lemmatize(word)) for word in review if ((word not in my_stopwords) and len(word) > 1)]
    review = [word for word in review if len(word) > 1]
    review=' '.join(review)
    return review
```

```
In [8]: df["processed_review"] = df.text.apply(clean_review)
df.head()
```

Out[8]:

	airline_sentiment	airline	text	processed_review
0	1	2	@VirginAmerica What @dhepburn said.	what said
1	1	2	@VirginAmerica plus you've added commercials t...	plu ve ad commerci experi tacki
2	1	2	@VirginAmerica I didn't today... Must mean I n...	did not today must mean need take anoth trip
3	0	2	@VirginAmerica it's really aggressive to blast...	realli aggress blast obnoxio not your guest fac...
4	0	2	@VirginAmerica and it's a really big bad thing...	and realli big bad thing about

```
In [9]: X = df.processed_review
y = df.airline_sentiment
```

```
In [10]: df.processed_review.sample(5).tolist()
```

```
Out[10]: ['worst nonrefund first class ticket oh becaus when select global fc their syst
em auto select economi upgrad',
'did start claim but week unrealist realli suppos go long with out car seat ch
ild ridicul',
'just promot product all problem with southwest and recommend noneoth than bes
t http co tfanxbh cf',
'fli from love austin now most not news',
'realli aggress blast obnoxio not your guest face amp littl recours']
```

```
In [11]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25, shuffle=True,
```

```
In [12]: from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer  
  
         cv=TfidfVectorizer(max_features=1200)  
  
         X_train=cv.fit_transform(X_train)  
         X_test=cv.transform(X_test)
```

```
In [13]: from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier

from keras.layers import InputLayer, Dense, Activation, Dropout
from keras.models import Sequential

model = Sequential()
model.add(Dense(1200, input_shape=X_train.shape[1:], activation="sigmoid"))

model.add(Dense(1, activation="sigmoid"))
model.compile(optimizer="adam", loss="binary_crossentropy", metrics=["accuracy"])
model.fit(X_train, y_train, epochs=25)
```

Using TensorFlow backend.

WARNING:tensorflow:From C:\Users\Admin\Anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py:74: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From C:\Users\Admin\Anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py:517: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From C:\Users\Admin\Anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py:4138: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.

WARNING:tensorflow:From C:\Users\Admin\Anaconda3\lib\site-packages\keras\optimizers.py:790: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From C:\Users\Admin\Anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py:3376: The name tf.log is deprecated. Please use tf.math.log instead.

WARNING:tensorflow:From C:\Users\Admin\Anaconda3\lib\site-packages\tensorflow\python\ops\nn_impl.py:180: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

WARNING:tensorflow:From C:\Users\Admin\Anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py:986: The name tf.assign_add is deprecated. Please use tf.compat.v1.assign_add instead.

Epoch 1/25

385/385 [=====] - 1s 3ms/step - loss: 0.6788 - acc: 0.5818

Epoch 2/25

385/385 [=====] - 1s 2ms/step - loss: 0.6458 - acc: 0.6286

Epoch 3/25

385/385 [=====] - 1s 2ms/step - loss: 0.6542 - acc: 0.5584

Epoch 4/25

```
385/385 [=====] - 1s 2ms/step - loss: 0.5904 - acc:
0.6494
Epoch 5/25
385/385 [=====] - 1s 2ms/step - loss: 0.5213 - acc:
0.7403
Epoch 6/25
385/385 [=====] - 1s 2ms/step - loss: 0.5095 - acc:
0.9117
Epoch 7/25
385/385 [=====] - 1s 2ms/step - loss: 0.5107 - acc:
0.6727
Epoch 8/25
385/385 [=====] - 1s 2ms/step - loss: 0.4792 - acc:
0.8519
Epoch 9/25
385/385 [=====] - 1s 2ms/step - loss: 0.4176 - acc:
0.7974
Epoch 10/25
385/385 [=====] - 1s 2ms/step - loss: 0.3851 - acc:
0.9273
Epoch 11/25
385/385 [=====] - 1s 2ms/step - loss: 0.3615 - acc:
0.9481
Epoch 12/25
385/385 [=====] - 1s 2ms/step - loss: 0.3540 - acc:
0.9636
Epoch 13/25
385/385 [=====] - 1s 2ms/step - loss: 0.3224 - acc:
0.9844
Epoch 14/25
385/385 [=====] - 1s 2ms/step - loss: 0.3042 - acc:
0.9506A: 0s - loss: 0.3246 - a
Epoch 15/25
385/385 [=====] - 1s 2ms/step - loss: 0.3098 - acc:
0.9714
Epoch 16/25
385/385 [=====] - 1s 2ms/step - loss: 0.2731 - acc:
0.9610
Epoch 17/25
385/385 [=====] - 1s 2ms/step - loss: 0.2496 - acc:
0.9792
Epoch 18/25
385/385 [=====] - 1s 2ms/step - loss: 0.2449 - acc:
0.9558
Epoch 19/25
385/385 [=====] - 1s 2ms/step - loss: 0.2282 - acc:
0.9844
Epoch 20/25
385/385 [=====] - 1s 2ms/step - loss: 0.2131 - acc:
0.9740
Epoch 21/25
385/385 [=====] - 1s 2ms/step - loss: 0.2189 - acc:
0.9766
Epoch 22/25
385/385 [=====] - 1s 2ms/step - loss: 0.1981 - acc:
0.9688
Epoch 23/25
```

```

385/385 [=====] - 1s 2ms/step - loss: 0.1867 - acc: 0.9766
Epoch 24/25
385/385 [=====] - 1s 2ms/step - loss: 0.1593 - acc: 0.9948
A: 0s - loss: 0.1592 - acc: 0.994
Epoch 25/25
385/385 [=====] - 1s 2ms/step - loss: 0.1490 - acc: 0.9948


```

Out[13]: <keras.callbacks.History at 0x28b9c3481c8>

In [14]: `y_pred=model.predict(X_test)`

In [15]: `plt.plot(y_pred)`

Out[15]: [`<matplotlib.lines.Line2D at 0x28b9c774848>`]



In [16]: `y_pred_activated = y_pred >= 0.5`

In [17]: `y_pred_activated.shape`

Out[17]: (129, 1)

In [18]: `y_test.shape`

Out[18]: (129,)

In [19]: `from sklearn.metrics import confusion_matrix`
`cm=confusion_matrix(y_test,y_pred_activated)`
`cm`

Out[19]: array([[31, 12],
[19, 67]], dtype=int64)

```
In [20]: from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
print(classification_report(y_test, y_pred_activated))
```

	precision	recall	f1-score	support
0	0.62	0.72	0.67	43
1	0.85	0.78	0.81	86
accuracy			0.76	129
macro avg	0.73	0.75	0.74	129
weighted avg	0.77	0.76	0.76	129

```
In [21]: print(accuracy_score(y_test, y_pred_activated))
```

0.7596899224806202

```
In [ ]:
```